

Samarth Gugnani

Software Engineer, Revenue and GTM Systems — six years building revenue infrastructure from zero

I build the systems revenue teams run on: account data, enrichment, qualification and routing, sequencing and deliverability, and the internal tooling that replaces manual work for sales and marketing. I wire the GTM stack against its APIs rather than a no-code canvas, and I build the guardrails that decide where an agent runs alone and where a human signs off.

Remote | Six years remote for distributed teams | **Overlap hours to suit: US, EU or APAC**

samarth.gugnani30@gmail.com | linkedin.com/in/samarth30 | github.com/samarth30

HIGHLIGHTS

- Revenue infrastructure at scale.** Built the credits and API-key billing platform that thousands of hosted AI agents ran on at **Eliza Labs** (15,000+ GitHub stars), on Cloudflare Workers, through a **\$20M round**.
- Signal and account data.** Built a hiring-signal crawler over five applicant-tracking APIs (Greenhouse, Lever, Ashby, Workable, Teamtailor): **351 companies verified live at the source** and **456 contacts resolved** across 219 of them, each row stamped with its evidence.
- LLM pipelines in production.** Thirteen chained prompts turning a prospect's own site into a personalized outbound sequence in **under ten minutes**, across **5,000+ accounts** sourced from live ad activity, with abstention when a claim cannot be traced to its source.
- Entity resolution.** Reverse-engineered TikTok's private API and resolved the same creator across **five platforms** that agreed on nothing into one schema, cutting data-layer response times **60%**.

PROFESSIONAL SUMMARY

Software engineer with six years building backend and data infrastructure, most of it as one of the first engineers at venture-backed startups, building from zero and working remotely for distributed teams, and most recently on revenue and go-to-market systems. This year I built a full GTM data stack from scratch: a hiring-signal crawler across five applicant-tracking APIs, a technology-detection scanner, enrichment and verification, and the sequencing infrastructure on top. I learned where an automated agent cannot be trusted the hard way. My own extractor once read a truncated job description and almost reported the wrong CRM into a live campaign. Now the extractors abstain when confidence is low and every row carries its evidence. Deciding what an agent settles alone and what a human signs off is the part of this work I care most about.

TECHNICAL SKILLS

- Languages, Backend and Data:** Python, TypeScript, JavaScript, Node.js, SQL, REST API design, webhooks, microservices, PostgreSQL, MySQL, MongoDB, Redis, DuckDB, Parquet, data modeling, ETL and incremental pipelines, entity resolution, deduplication, data warehouse querying, analytics dashboards
- GTM and RevOps Systems:** Clay, Apollo, Instantly, Smartlead, n8n, Zapier, Make, HubSpot, CRM and account data modeling, lead qualification and routing, waterfall enrichment, email verification, deliverability, signal detection (hiring, funding, technology install), GDPR and CAN-SPAM aware sending, internal tools for operators
- AI and Automation:** Anthropic Claude API, OpenAI API, Claude Code, Cursor, MCP integrations, LLM workflows, prompt engineering, multi-agent pipelines, adversarial verification, embeddings and vector search, signal extraction and classification from unstructured text
- Crawling, Frontend and Infrastructure:** async HTTP crawlers, Playwright, web scraping, DNS and certificate-transparency analysis, response-header fingerprinting, rate limiting, private-API reverse engineering, Next.js, React, Tailwind, Cloudflare Workers, AWS Lambda, Vercel, Railway, Docker, Git, CI/CD

PROFESSIONAL EXPERIENCE

Software Engineer, GTM Systems | Independent Consulting (self-directed build, shipped to production)

Remote | Apr 2026 to Present

Built a revenue-systems stack from scratch: *signal layer, account data model, enrichment and verification, sequencing, deliverability and operator tooling.*

- Signal layer.** Hiring-signal crawler pulling requisitions directly from applicant-tracking APIs and extracting the tools named in each. Cached indexes keep listing filled roles as open for weeks, so every posting is re-checked at the source before use, and each row is stamped with its source API and timestamp.
- Technology detection from website code.** Scanner that determines which software a company runs by reading DNS records, HTTP response headers, cookie names and the hosts its pages load resources from, pulling out a per-customer ID such as a HubSpot portal ID or a GA4 measurement ID as proof. A vendor merely named in page text is not an install and is routed to a rejected list.
- Sequencing and deliverability infrastructure.** Wrote roughly 1,500 lines of Python against a sending platform's API covering campaign push, rate caps per inbox and per domain, send scheduling in each recipient's own timezone, and a suppression registry. Where the documentation was silent I probed the API to determine how it actually behaved.
- LLM extraction pipeline.** Thirteen chained prompts converting a rendered screenshot of a prospect's site into a personalized outbound sequence. Any claim not traceable to the source page is dropped before send.
- Enrichment.** Enriched the account and contact tables through the **Clay** CLI, waterfall enrichment then email verification, and pushed the cleaned segments into **Instantly** campaigns. Rows that cannot be verified are quarantined with a reason rather than sent.
- Built with coding agents.** Built this stack with **Claude Code** and **Cursor** as daily tools rather than experiments: custom skills, subagents and MCP integrations for the work that repeats, and a Grok bot for the rest. What an agent is allowed to settle without me is decided up front.
- Internal tooling.** Built the operator-facing review and tracking surfaces where a human accepts, rejects or corrects what the automation proposed before it ships, in **Next.js, React and TypeScript**, live in production.

Senior Software Engineer | Eliza Labs

San Francisco Bay Area (Remote) | Nov 2024 to Mar 2026

One of the first engineers, reporting directly to the CTO. Core maintainer of an open-source AI-agent operating system with 15,000+ GitHub stars and 100+ contributors.

- Built the credits and API-key billing platform** that thousands of hosted AI agents ran on, at the edge on Cloudflare Workers, covering key issuance, usage metering, accounting and quota enforcement. These were long-running agents in production rather than scripts, so the platform had to hold quotas and accounting correct under concurrency for thousands of them at once.
- Built a **plugin autocoder** that generates boilerplate for a new plugin from a spec, cutting setup time by **70%**. Reviewed and merged work from **100+ contributors** and mentored interns, fully async across a 12-hour time difference.
- Joined as one of the first few engineers and stayed through a **\$20M funding round**, as the project became one of the largest open-source AI-agent ecosystems and its token peaked at a **\$2.5B market cap**.

Full-Engineer | Camp Network

Berkley, Michigan (Remote) | Jun 2023 to Oct 2024

One of the first few engineers. Owned the data and AI layer behind a creator-analytics platform from pre-seed to Series A, across **\$30M raised**.

- Reverse-engineered TikTok's API** and wrote scrapers for data its public endpoints did not expose, then solved **entity resolution** across TikTok, YouTube, Twitter, Instagram and Spotify, matching the same creator across five sources that agreed on nothing into one schema. Cut response times across the data layer by **60%**.
- Shipped an AI chatbot answering metrics questions on demand, lifting satisfaction **30%** and cutting support tickets **40%**.

Full-Stack Engineer | Shield (a16z-backed), Freelance

Remote | Mar 2023 to Jun 2023

- Zero-to-one build shipped in four months, backend through UI, as the engineer on the build rather than one of a team. Designed the data model and APIs, then shipped against live user feedback until the core use case held. Lifted customer acquisition **25%**.

Chief Technology Officer | SignAssist

Remote | Jun 2022 to Dec 2022

Ran engineering for a seven-person company and took the product through a full pivot.

- Shipped a Chrome extension that made transactions readable and safe for end users, cutting fraud **45%** and lifting engagement **20%**.
- Built the **API-as-a-service business that replaced it** and sold it into B2B customers, **landing Frontier.xyz as a client**. Ran 50+ customer interviews, then repositioned the product and business model on what came back, lifting customer acquisition **20%** in three months.

Founder | Intelli Chains

Remote | Aug 2021 to Jun 2022

- Ran an engineering services company, **delivering for 15+ international clients in under a year** across marketplaces, subscription products and lending applications. Designed the APIs and distributed systems behind them, built **real-time indexing bots over high-throughput event streams**, and managed interns.

Earlier Experience

- Full-Stack Engineer, Crypton Network** (Oct 2020 to Sep 2021), second employee at an early-stage fintech; built the core transaction systems and test suite as platform assets scaled to **\$10M market cap**. **Blockchain Lead, YFDAI Finance** (part-time, 2020 to 2021). **Full-Stack Engineer, Unbox Innovations** (internship, 2020).

EDUCATION AND OPEN SOURCE

- Chitkara University**, B.Tech, Computer Science. **Google Summer of Code** (Amahi, 2020), merged PRs on system reliability and built non-admin user login.